

Density estimation

Computational Statistics

Adam Gorm Hoffmann

Goal

Compute kernel density estimate

$$\hat{f}_h(x) = \frac{1}{hn} \sum_{j=1}^n K\left(\frac{x - x_j}{h}\right)$$

where

$$K(x) = \frac{3}{4}(1 - x^2)1_{[-1,1]}(x)$$

is the Epanechnikov kernel.

- ▶ First: How and where to compute the kernel density estimate?
 - ▶ Basic implementation.
 - ▶ Different binning implementations (R and Rcpp).
- ▶ Then (no time...): How to choose a reasonable bandwidth h ?
 - ▶ k -fold cross validation.
 - ▶ AMISE plug-in.

Where to estimate density?

An extreme case would be to estimate the density in all x_1, \dots, x_n :

$$\hat{f}_h(x_i) = \frac{1}{hn} \sum_{j=1}^n K\left(\frac{x_i - x_j}{h}\right)$$

which requires n^2 kernel evaluations ($\frac{1}{2}(n-1)^2 + 1$ when using symmetry of K).

Instead we only estimate it in grid points g_1, \dots, g_m :

$$\hat{f}_h(g_i) = \frac{1}{hn} \sum_{j=1}^n K\left(\frac{g_i - x_j}{h}\right)$$

which requires (half of) mn kernel evaluations.

C++ functions and R wrapper

```
1 double
2 ep_density_single_cpp(double x, const NumericVector& x_obs, double h)
3 {
4     double result = 0;
5     unsigned long long n = x_obs.length();
6     for (unsigned long long i = 0; i < n; i++)
7         result += ep_kernel_single_cpp((x - x_obs[i]) / h);
8     return result / (h * n);
9 }
10
11 // [[Rcpp::export]]
12 NumericVector
13 ep_density_cpp(const NumericVector& x_obs, double h, const NumericVector& grid)
14 {
15     NumericVector density_estimates(grid.length());
16     for (int i = 0; i < grid.length(); i++)
17         density_estimates[i] = ep_density_single_cpp(grid[i], x_obs, h);
18     return density_estimates;
19 }
```

R wrapper:

```
1 ep_density_cpp_wrap <- function(x_obs, h, m = 512)
2 {
3     grid <- seq(min(x_obs) - 3 * h, max(x_obs) + 3 * h, length.out = m)
4     list(
5         x = grid,
6         y = ep_density_cpp(x_obs, h, grid)
7     )
8 }
```

Binning: further reducing kernel evaluations

Assign each grid point g_j the grid count

$$w_j = \sum_{i=1}^n 1(g_j \text{ is the grid point closest to } x_i)$$

and use the kernel density estimate

$$\hat{f}_h(g_i) = \frac{1}{hn} \sum_{j=1}^m w_j \cdot K\left(\frac{g_i - g_j}{h}\right)$$

which requires only (half of) m^2 kernel evaluations.

Less kernel evaluations at cost of accuracy.

Even fewer evaluations with equidistant grid

Note that when g_1, \dots, g_m are equidistant as

$$g_1, g_1 + \delta, g_1 + 2\delta, \dots, g_1 + (m-1)\delta$$

then

$$|g_i - g_j| \in \{0, \delta, \dots, (m-1)\delta\}, \quad \text{for all } i, j = 1, \dots, m$$

so we only have to do m kernel evaluations

$$K\left(\frac{0}{h}\right), K\left(\frac{\delta}{h}\right), \dots, K\left(\frac{(m-1)\delta}{h}\right)$$

when we also use that K is symmetric.

And even fewer for kernels where $K(y) = 0$ for $|y| \geq 1$

Note that

$$K\left(\frac{i\delta}{h}\right) = 0, \quad \text{for } i \geq \frac{h}{\delta}$$

so we only have to calculate

$$K(0), K\left(\frac{\delta}{h}\right), \dots, K\left(\text{floor}\left(\frac{h}{\delta}\right) \frac{\delta}{h}\right)$$

which amounts to $\text{ceiling}\left(\frac{h}{\delta}\right)$ kernel evaluations.

Example: Standard Gaussian

We choose the following reasonable parameters:

- ▶ $n = 10^5$ observations.
- ▶ $m = 512$ grid points in $[-3\sigma, 3\sigma] = [-3, 3]$.
- ▶ $\delta = \frac{3 - (-3)}{512} \approx 0.01$.
- ▶ $h = 1.06\sigma n^{-1/5} \approx 0.1$.

Number of kernel evaluations	Method
$n^2 = 10^{10}$	evaluating in all observations
$mn = 512 \cdot 10^5 = 5.12 \cdot 10^7$	evaluating in grid
$m^2 \approx 2.6 \cdot 10^5$	naive binning
$m = 512$	naive equidistant binning
$\text{ceiling}\left(\frac{h}{\delta}\right) = 10$	smart equidistant binning

We save a lot of kernel evaluations! But we still need some iterations to sum the weighted kernel values.

R implementation of weights

```
1 bin_weights <- function(x, grid_lower, grid_length, grid_diff)
2 {
3   w <- numeric(grid_length)
4   for(i in seq_along(x)) {
5     closest_grid_i <- ceiling((x[i] - grid_lower) / grid_diff + 0.5)
6     w[closest_grid_i] <- w[closest_grid_i] + 1
7   }
8   w
9 }
```

► Very easy to translate to C++.

Rcpp implementation of weights

```
1 NumericVector
2 bin_weights_cpp(const NumericVector& x, double grid_lower,
3                 int grid_length, double grid_diff)
4 {
5     double closest_grid_i;
6     NumericVector w(grid_length);
7     unsigned long long n_x = x.length();
8
9     for (unsigned long long i = 0; i < n_x; i++) {
10         closest_grid_i = floor((x[i] - grid_lower) / grid_diff + 0.5);
11         w[closest_grid_i] += 1;
12     }
13
14     return w;
15 }
```

R implementation of binning

```
1  ep_density_binning <- function(x, h, grid_length = 512)
2  {
3      grid_lower <- min(x) - 3 * h
4      grid_upper <- max(x) + 3 * h
5      grid <- seq(grid_lower, grid_upper, length.out = grid_length)
6      grid_diff <- grid[2] - grid[1]
7      max_nonzero_i <- floor(h / grid_diff)
8      kernel_evals <- ep_kernel_cpp((grid[1:(max_nonzero_i + 1)] -
9                                     grid_lower) / h)
10     kernel_vec <- c(rev(kernel_evals[-1]), kernel_evals)
11     weights <- c(rep(0, max_nonzero_i),
12                  bin_weights_cpp(x, grid_lower, grid_length, grid_diff),
13                  rep(0, max_nonzero_i))
14     y <- numeric(grid_length)
15     for (i in (1 + max_nonzero_i):(grid_length + max_nonzero_i))
16         y[i - max_nonzero_i] <- sum(weights[(i - max_nonzero_i):
17                                             (i + max_nonzero_i)] *
18                                     kernel_vec)
19
20     list(x = grid, y = y / (h * length(x)))
21 }
```

R implementation with Toeplitz trick

The kernel is evaluated in all grid points, but all computations are vectorized.

```
1 ep_density_binning_toeplitz <- function(x, h, grid_length = 512)
2 {
3     n <- length(x)
4     grid_lower <- min(x) - 3 * h
5     grid_upper <- max(x) + 3 * h
6     grid <- seq(grid_lower, grid_upper, length.out = grid_length)
7     grid_diff <- grid[2] - grid[1]
8     kernel_evals <- ep_kernel_cpp((grid - grid_lower) / h)
9     weights <- bin_weights_cpp(x, grid_lower, grid_length, grid_diff)
10    list(
11        x = grid,
12        y = colSums(weights * toeplitz(kernel_evals)) / (h * length(x))
13    )
14 }
```

Problems with profiling the implementations

- ▶ They are too fast to be profiled.
- ▶ When profiling a replication of them (so they take longer), the results vary wildly.

Solution

- ▶ Benchmark each line manually with `bench::mark`
- ▶ I do it for 10^8 observations x .
- ▶ Possible package idea? Seems useful!

Benchmarking: `bin_weights`

R implementation: 16.3s

Rcpp implementation: 450ms

As expected, since we didn't vectorize anything in R.

Benchmarking: min and max are slow! So is range.

Expression	Median time (ms)
<code>c(min(x), max(x))</code>	259
<code>range(x)</code>	715
<code>range_cpp(x)</code>	97

Solution: I wrote my own fast range copy in C++.

```
1 NumericVector
2 range_cpp(const NumericVector& x)
3 {
4     unsigned long long n_x = x.length();
5     NumericVector range(2);
6     range[0] = x[0];
7     range[1] = x[0];
8     for (unsigned long long i = 1; i < n_x; i++) {
9         if (x[i] < range[0])
10             range[0] = x[i];
11         else if (x[i] > range[1])
12             range[1] = x[i];
13     }
14     return range;
15 }
```

Benchmark-profiling the smart implementation

Kernel evaluations: 0.002ms.

bin_weights_cpp: 458ms.

The loop: 5ms.

Bottlenecks: bin_weights_cpp, min and max.

- ▶ min and max can't easily be improved more than range_cpp.
 - ▶ Possible solution is to let the user supply end-points of grid manually, if they have prior knowledge of a reasonable grid.
- ▶ bin_weights_cpp seems to be optimized in C++.
- ▶ Note: This is for very a very large data set (10^8 obs.)
- ▶ When the number of observations is not much larger than the grid length, the loop and kernel evaluations matter more (so it still makes sense to optimize them in C++, as we will see).

Benchmark-profiling Toeplitz implementation

Kernel evaluations: 0.005ms.

bin_weights_cpp: 458ms.

toeplitz(kernel_evals): 6ms.

Multiplying with weights: 1ms.

colSums: 0.2ms.

- ▶ weights, min, and max are still the largest bottlenecks.
- ▶ The Toeplitz implementation is a few ms slower than the smart R implementation – the grid length is so little compared to the length of x, that doing 512 kernel evaluations instead of 4 doesn't really matter.

Conclusions from benchmark-profiling

- ▶ Lines that scale with the size of `x` are the bottlenecks (`bin_weights`, `min`, `max`).
- ▶ We have reduced their runtime significantly with C++ implementations, but they are still the bottlenecks.
- ▶ We will now do a complete C++ implementation, and see that it still saves time for smaller sample sizes.

Rcpp implementation: highest level

```
1 NumericVector
2 ep_density_binning_cpp(const NumericVector& x, double h,
3                        const NumericVector& grid)
4 {
5     int i;
6     int n_x = x.length();
7     int grid_length = grid.length();
8     double grid_lower = grid[0];
9     double grid_diff = grid[1] - grid[0];
10    NumericVector y(grid_length);
11    NumericVector weights = bin_weights_cpp(x, grid_lower,
12                                           grid_length, grid_diff);
13    int max_nonzero_i = floor(h / grid_diff);
14    double kernel_evals[max_nonzero_i + 1];
15    kernel_evals[0] = 0.75; // Ep kernel in 0.
16    for (i = 1; i <= max_nonzero_i; i++)
17        kernel_evals[i] = ep_kernel_single_cpp((i * grid_diff) / h);
18    for (i = 0; i < grid_length; i++)
19        y[i] = weighted_kernel_sum(i, max_nonzero_i, grid_length,
20                                   weights, kernel_evals) / (h * n_x);
21    return y;
22 }
```


Rcpp implementation: weighted kernel sum

```
1  double
2  weighted_kernel_sum(double i, double max_nonzero_i, double grid_length,
3                      const NumericVector& weights,
4                      const double *kernel_evals)
5  {
6      int below, above;
7      double res = weights[i] * *kernel_evals;
8      for (int j = 1; j <= max_nonzero_i; j++) {
9          below = i - j;
10         above = i + j;
11         if (below >= 0)
12             res += weights[below] * *(kernel_evals + j);
13         if (above < grid_length)
14             res += weights[above] * *(kernel_evals + j);
15     }
16
17     return res;
18 }
```

Rcpp R wrapper

```
1 ep_density_binning_cpp_wrap <- function(x, h, grid_length = 512)
2 {
3     grid <- seq(min(x) - 3 * h, max(x) + 3 * h, length.out = grid_length)
4     list(
5         x = grid,
6         y = ep_density_binning_cpp(x, h, grid)
7     )
8 }
```

Benchmarking on Gaussian mixture



- As expected, C++ doesn't make much difference for very large samples.

Finally: my own density copy

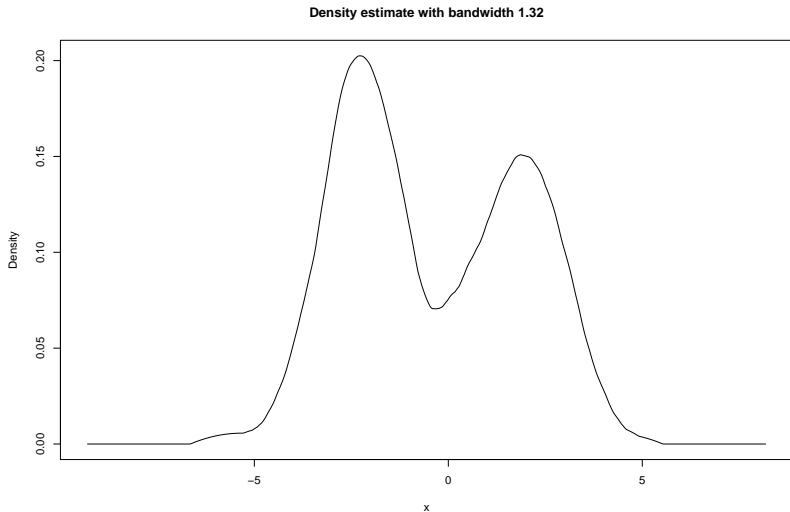
```
1  my_density <- function(x, bw = "plug-in", binning = TRUE, ...)  
2  {  
3    if (bw == "plug-in")  
4      h <- bandwidth_plug_in_precalculated(x)  
5    else if (bw == "cv")  
6      h <- bandwidth_cv(x, ...)  
7    else if (bw == "silverman")  
8      h <- bandwidth_silverman(x)  
9    else if (is.function(bw))  
10     h <- bw(x, ...)  
11    else if (is.numeric(bw))  
12     h <- bw  
13    else  
14     stop("Invalid bw.")  
15  
16    if (binning)  
17     xy <- ep_density_binning_cpp_wrap(x, h, ...)  
18    else  
19     xy <- ep_density_cpp_wrap(x, h, ...)  
20  
21    structure(  
22      list(  
23        x = xy$x,  
24        y = xy$y,  
25        bw = h  
26      ),  
27      class = "my_density"  
28    )  
29  }
```

Plot method


```
1 plot.my_density <- function(my_density_object, ...)  
2   plot(my_density_object$x, my_density_object$y,  
3       type = "l", xlab = "x", ylab = "Density",  
4       main = paste("Density estimate with bandwidth",  
5                     round(my_density_object$bw, digits = 2)),  
6       ...)
```


Testing on Gaussian mixture: plot method.


```
1 plot(my_density(x, bw = "plug-in", binning = FALSE))
```




my_density and density with same bandwidth




Differences between density and non-binning imp.



Differences between density and binning imp.




Differences between binning and non-binning imp.



Bandwidth selection with k -fold cross-validation and plug-in

- ▶ k -fold cross validation
 - ▶ Improve speed by caching complements of folds I^{-i} .
- ▶ Plug-in
 - ▶ Improve speed with Rcpp implementation.



k -fold cross-validation

Select bandwidth

$$\hat{h}_{\text{CV}(k)} = \arg \max_h \sum_{i=1}^n \log(\hat{f}_h^{-i}).$$

where

$$\hat{f}_h^{-i} = \frac{1}{hn_i} \sum_{j \in I^{-i}} K\left(\frac{x_i - x_j}{h}\right)$$

is the kernel density estimate in x_i based on the $k - 1$ folds that don't contain i .

Highest level of implementation

```
1  ## Selects h by k-fold cross validation
2  bandwidth_cv <- function(x_obs, k)
3  {
4      optimize(
5          f = get_cv_objective(x_obs, k),
6          interval = c(0, max(x_obs) - min(x_obs)),
7          maximum = TRUE
8      )$maximum
9  }
```


Naive objective function implementation

```
1  get_cv_objective <- function(x_obs, k)
2  {
3      force_all_args()
4      n <- length(x_obs)
5      ## All the indices with a 3 is I_3, and so on.
6      fold_indices <- rep(1:k, times = ceiling(n / k))[sample(n)]
7
8      function(h)
9      {
10         dens_estimates <- numeric(n)
11         for (i in 1:n) {
12             ## All x-values that are not in the same fold as the i'th
13             ## observation ( $I^{-i}$ ).
14             x_other_folds <- x_obs[fold_indices != fold_indices[i]]
15             dens_estimates[i] <- ep_density_single(x_obs[i],
16                                                     x_other_folds,
17                                                     h)
18         }
19         sum(log(dens_estimates))
20     }
21 }
22
23 ## Forces evaluation of all arguments
24 force_all_args <- function()
25     as.list(parent.frame())
```

Profiling bandwidth_cv

10-fold cross validation on standard Gaussian data.

```
1 library(profvis)
2 source("implementation.R", keep.source = TRUE)
3
4 n_sim <- 10000
5 x <- rnorm(n_sim)
6
7 profvis(bandwidth_cv(x, 10))
```

Profiling bandwidth_cv: Objective function

```
get_cv_objective_naive <- function(x_obs, k)
{
  n <- length(x_obs)
  ## All the indices with a 3 is I_3, and so on.
  fold_indices <- rep(1:k, times = ceiling(n / k))[sample(n)]

  ## Function to maximize
  function(h)
  {
    dens_estimates <- numeric(n)
    for (i in 1:n) {
      ## All x-values that are not in the same fold as the i'th
      ## observation.
      x_other_folds <- x_obs[fold_indices != fold_indices[i]]
      dens_estimates[i] <- ep_density_single(x_obs[i], x_other_folds, h)
    }
    ## Return log-likelihood.
    sum(log(dens_estimates))
  }
}
```

8510

21970

- ▶ Problem: We recalculate I^{-i} for each element in a given fold, and indexing is slow in R.
- ▶ Idea: Cache I^{-i} the first time we meet an element from a given fold, and reuse it for later elements.

Profiling bandwidth_cv: ep_kernel

```
ep_kernel <- function(x)
{
  nonzero_index <- abs(x) <= 1
  result <- numeric(length(x))
  result[nonzero_index] <- (1 - x[nonzero_index]^2) * 3/4
  result
}
```



► Idea: Implement in C++

(R implementation $(\text{abs}(x) \leq 1) * (1 - x^2) * 0.75$ is faster in some cases, but let us just go straight to C++).

Objective function implementation with caching

```
1  get_cv_objective <- function(x_obs, k)
2  {
3      force_all_args()
4      n <- length(x_obs)
5      ## All the indices with a 3 is I_3, and so on.
6      fold_indices <- rep(1:k, times = ceiling(n / k))[sample(n)]
7      other_folds_list <- vector("list", k)
8      function(h)
9      {
10         dens_estimates <- numeric(n)
11         for (i in 1:n) {
12             ## All x-values that are not in the same fold as the i'th
13             ## observation ( $I^{-i}$ ).
14             x_other_folds <- other_folds_list[[fold_indices[i]]]
15             if (is.null(x_other_folds)) {
16                 x_other_folds <- x_obs[fold_indices != fold_indices[i]]
17                 other_folds_list[[fold_indices[i]]] <- x_other_folds
18             }
19             dens_estimates[i] <- ep_density_single(x_obs[i],
20                                                     x_other_folds,
21                                                     h)
22         }
23         sum(log(dens_estimates))
24     }
25 }
```

Rcpp implementation of ep_kernel

```
1  double
2  ep_kernel_single_cpp(double x)
3  {
4      if (-1 < x && x < 1)
5          return (1 - square(x)) * 0.75;
6      else
7          return 0;
8  }
9
10 // [[Rcpp::export]]
11 NumericVector
12 ep_kernel_cpp(const NumericVector& x)
13 {
14     unsigned long long n = x.length();
15     NumericVector y(n);
16     for (unsigned long long i = 0; i < n; i++)
17         y[i] = ep_kernel_single_cpp(x[i]);
18     return y;
19 }
```

Profiling again

```
get_cv_objective <- function(x_obs, k)
{
  n <- length(x_obs)
  ## All the indices with a 3 is I_3, and so on.
  fold_indices <- rep(1:k, times = ceiling(n / k))[sample(n)]
  other_folds_list <- vector("list", k)

  function(h)
  {
    dens_estimates <- numeric(n)
    for (i in 1:n) {
      ## All x-values that are not in the same fold as the i'th
      ## observation (I^{i}).
      x_other_folds <- other_folds_list[[fold_indices[i]]]
      if (is.null(x_other_folds)) {
        x_other_folds <- x_obs[fold_indices != fold_indices[i]]
        other_folds_list[[fold_indices[i]]] <- x_other_folds
      }
      dens_estimates[i] <- ep_density_single(x_obs[i],
                                             x_other_folds,
                                             h)
    }
    sum(log(dens_estimates))
  }
}
```

50


13840

Finding I^{-i} is no longer a bottleneck.


Benchmarking the ep_kernel implementations

```
1  vectors <- lapply(10^(1:8), rnorm)
2  function_names <- c("ep_kernel", "ep_kernel_cpp")
3  call_grid <- expand.grid(funs = function_names,
4                          vector_indices = 1:length(vectors))
5  calls <- paste0(call_grid$funs, "(vectors[[" ,
6                  call_grid$vector_indices, "]])")
7  expression_list <- lapply(
8    calls,
9    function(call_str) parse(text = call_str)[[1]]
10 )
11 benchmark <- bench::mark(exprs = expression_list,
12                           min_iterations = 5,
13                           check = FALSE)
```


Rcpp implementation more than halves runtime



Runtime scales linearly with sample size



CV problem with outliers

$$\min_{j \in I^{-i}} (|x_i - x_j|) = d \Rightarrow \hat{f}_h^{-i} = 0 \text{ for all } h \leq d$$

$$\Rightarrow \log(\hat{f}_h^{-i}) = -\infty$$

$$\Rightarrow \sum_{i=1}^n \log(\hat{f}_h^{-i}) = -\infty$$

$$\Rightarrow \hat{h}_{CV(k)} = \operatorname{argmax}_h \sum_{i=1}^n \log(\hat{f}_h^{-i}) > d.$$

Hence

$$\hat{h}_{CV(k)} > \max_{i=1, \dots, n} \min_{j \in I^{-i}} (|x_i - x_j|).$$

One x_i far from all points in the other folds I^{-i} leads to a large $\hat{h}_{CV(k)}$ (oversmoothing).

Other approach: Minimize AMISE

- The Asymptotic Mean Integrated Squared Error (that is, the asymptotically dominating terms of the integral of the MSE) is minimized by

$$h_n = \left(\frac{\|K\|_2^2}{\|f_0''\|_2^2 \sigma_K^4} \right)^{1/5} n^{-1/5}.$$

- h_n is thus a reasonable choice of bandwidth for kernel density estimation... if we knew $\|f_0''\|_2^2$.

Plug-in estimate


- Use Silverman's rule of thumb to choose initial bandwidth $r = 0.9\tilde{\sigma}n^{-1/5}$, where $\tilde{\sigma} = \min\{\hat{\sigma}, \text{IQR}/1.34\}$.
- Estimate $\|f_0''\|_2^2$ using a Gaussian kernel H and bandwidth r :

$$\begin{aligned}\|\tilde{f}''\|_2^2 &= \frac{1}{n^2 r^6} \sum_{i=1}^n \sum_{j=1}^n \int H''\left(\frac{x - x_i}{r}\right) H''\left(\frac{x - x_j}{r}\right) dx \\ &= \frac{1}{n^2 (\sqrt{2}r)^5} \sum_{i=1}^n \sum_{j=1}^n \phi^{(4)}\left(\frac{X_i - X_j}{\sqrt{2}r}\right)\end{aligned}$$

(Bandwidth Selection: Classical or Plug-In?, Clive R. Loader, The Annals of Statistics, Apr., 1999, Vol. 27, No. 2, pp. 415-438)

- Use $\|\tilde{f}''\|_2^2$ to calculate \hat{h}_n .

Long story short: Rcpp wins



Final solution in Rcpp (highest level)

```
1  double
2  d2f0_two_norm_squared_estimate_cpp(const NumericVector& x, double r)
3  {
4      unsigned long long n = x.length();
5      double result = 0;
6      double rsqrt2 = r * sqrt(2);
7      for (unsigned long long i = 0; i < n - 1; i++)
8          for (unsigned long long j = i + 1; i < j && j < n; j++)
9              result += d4_std_gaussian_density_single_cpp((x[j] - x[i]) / rsqrt2);
10     return (2 * result + n * d4_std_gaussian_density_0 /
11            (square(n) * pow(rsqrt2, 5)));
12 }
```

Final solution in Rcpp (density)

```
1 double
2 d4_std_gaussian_density_single_cpp(double x)
3 {
4     double x_sq = square(x);
5     return exp(-x_sq/2) * ((x_sq - 6) * x_sq + 3) / sqrt2pi;
6 }
7
8 NumericVector
9 d4_std_gaussian_density_cpp(const NumericVector& x)
10 {
11     unsigned long long n = x.length();
12     NumericVector y(n);
13     for (unsigned long long i = 0; i < n; i++)
14         y[i] = d4_std_gaussian_density_single_cpp(x[i]);
15     return y;
16 }
```


Highest level of R implementation

```
1 bandwidth_plug_in <- function(x_obs)
2   (15 / (d2f0_two_norm_squared_estimate(x_obs) * length(x_obs)))^0.2
```

► We must implement `d2f0_two_norm_squared_estimate`.

► In R:

```
1 ## Estimate of the squared 2-norm of f0''.
2 d2f0_two_norm_squared_estimate <- function(x_obs,
3                                           r = bandwidth_silverman(x_obs),
4                                           diff_vec = diff_vec_expand.grid)
5   sum(d4_std_gaussian_density(diff_vec(x_obs) / (sqrt2 * r))) /
6     (length(x_obs)^2 * (sqrt2 * r)^5)
```

where `diff_vec` must return the vector of all pairwise differences of elements in `x_obs` (due to symmetry of the Epanechnikov kernel, we can actually get away with only calculating half of the differences).

Many ways to implement diff_vec...

```
1  ## Returns vector of all differences of pairs.
2  diff_vec_expand.grid <- function(x_obs)
3  {
4      x_grid <- expand.grid(Xi = x_obs, Xj = x_obs)
5      x_grid[, 1] - x_grid[, 2]
6  }
7
8  ## Same as above, but uses for loop
9  diff_vec_loop <- function(x_obs)
10 {
11     n <- length(x_obs)
12     res <- numeric(n^2)
13     count <- 1
14     for (i in 1:n) {
15         for (j in 1:n) {
16             res[count] <- x_obs[i] - x_obs[j]
17             count <- count + 1
18         }
19     }
20     res
21 }
22
23 ## Same as above, but uses "outer" instead of "expand.grid" and returns matrix.
24 diff_matrix_outer <- function(x_obs)
25     outer(x_obs, x_obs, FUN = "-")
26
27 ## Returns vector version of matrix result from above.
28 diff_vec_outer <- function(x_obs)
29     as.vector(diff_matrix_outer(x_obs))
```

... and a couple more.

```
1  ## One half of differences, since epa-kernel is symmetric
2  diff_half_combn <- function(x_obs)
3    combn(x_obs, 2, FUN = diff)
4
5  diff_half_outer <- function(x_obs)
6  {
7    diff_mat <- diff_matrix_outer(x_obs)
8    diff_mat[upper.tri(diff_mat)]
9  }
```

What else could be done?

- Implementation for general kernel (e.g., allow user to supply kernel S3 object, and the program has fall-back option of numerical integration if $\|\tilde{f}''\|_2^2$ is not implemented by the user).

Source on binning

Sections 5.2.1 and 5.3.2 of *Nonparametric Kernel Density Estimation and Its Computational Aspects* by Artur Gramacki (Springer 2018))

Old benchmark (without range_cpp)

